


The Redemption Optimization Appendix: Proof Sketch

Core Claim: Evil is necessary once (to enable redemption) but has no value beyond the first occurrence; repetition is suboptimal.

Evaluator $J = \Sigma(\alpha\Delta L + \gamma\Delta F + \mu M + \nu J_v) - \Sigma(\beta H + \kappa R)$. Credits are concave/typal; costs per rejection are positive. Redemption goods are event-valued (credited only when enacted).

Two Paths: (H0) Evil-free ($R=0$) yields no redemption credit unless redemption is valueless; (H1) Minimal-trigger ($R=1$) mints redemption, then structural confirmation closes the hazard while keeping freedom live.

Visual Aid: Objective J vs. Rejections under Typal Credits



Theorem (Sketch): For each incident class c with typal/concave credits and $\kappa > 0$, J is maximized at $n=1$. For $n > 1$, marginal credit ≤ 0 while marginal cost > 0 , so repeats are dominated. Therefore, one realized rejection suffices to ground mercy/justice; repetitions add harm without value.

Conclusion: Evil is structurally necessary once, never more. Confirmation closes hazards without coercion. This resolves the problem of evil under event-valuation and guides policy across domains.