


The Redemption Optimization: Simulation Results Summary

Objective: To test the Margan Optimization Paradox (MOP): evil is necessary once to ground redemption goods, but repetition adds no value.

Method: Event-valued ledger; concave/typal credits; linear costs; confirmation tests. Two stylized models are reported: Typal credits (all redemption credit upon the first rejection) and Concave credits (diminishing returns via square-root).


Results (Typal Credits)

Objective J peaks at $R = 1$ and declines thereafter: the first rejection grounds redemption; repeats add only cost.



Results (Concave Credits)

With diminishing-return credits ($10 \cdot \sqrt{R}$) and linear costs ($5 \cdot R$), the same qualitative pattern holds: J maximizes at a low R (typically at or near $R = 1$) and decreases as repetitions grow.



Conclusion

Across both models, simulations confirm minimal-trigger optimality: one realized rejection is sufficient to mint redemption in actuality; further rejections are dominated. This supports the policy doctrine of structural confirmation after the first breach.